

FDS PROJECT : DIABETES PREDICTION

Problem Statement : Predicting the likelihood of diabetes in individuals using health-related indicators, and identifying which factors contribute most to the diagnosis?

1. IMPORTING LIBRARIES

```
In [42]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
from scipy.stats import ttest_ind
```

2. LOAD AND EXPLORE THE DATASET

```
In [43]: df = pd.read_csv('pima-data.csv')

print(" 📄 First 5 records:")
print(df.head())
```

```
📄 First 5 records:
   num_preg  glucose_conc  diastolic_bp  thickness  insulin  bmi  diab_pred \
0         6          148           72         35         0  33.6      0.627
1         1           85           66         29         0  26.6      0.351
2         8          183           64          0         0  23.3      0.672
3         1           89           66         23         94  28.1      0.167
4         0          137           40         35        168  43.1      2.288

   age  skin  diabetes
0   50  1.3790      True
1   31  1.1426     False
2   32  0.0000      True
3   21  0.9062     False
4   33  1.3790      True
```

```
In [44]: print(df.tail())

   num_preg  glucose_conc  diastolic_bp  thickness  insulin  bmi \
763        10          101           76         48        180  32.9
764         2          122           70         27         0  36.8
765         5          121           72         23        112  26.2
766         1          126           60          0         0  30.1
767         1           93           70         31         0  30.4

   diab_pred  age  skin  diabetes
763    0.171   63  1.8912     False
764    0.340   27  1.0638     False
765    0.245   30  0.9062     False
766    0.349   47  0.0000      True
767    0.315   23  1.2214     False
```

3. HANDLING MISSING VALUES

```
In [45]: # Quick summary of missing values
print(df.isnull().sum())
```

```

num_preg      0
glucose_conc  0
diastolic_bp  0
thickness     0
insulin       0
bmi           0
diab_pred     0
age           0
skin          0
diabetes      0
dtype: int64

```

```
In [46]: df['age_group'] = pd.cut(df['age'], bins=[0, 30, 40, 50, 60, 100], labels=['<30', '30s', '40s', '50s', '60s', '70s', '80s', '90s', '100s'])
```

```
In [47]: print("\n 📐 Shape of dataset:", df.shape)
print("\n 📄 Column names:", df.columns.tolist())
print("\n 📊 Data types:\n", df.dtypes)
```

📐 Shape of dataset: (768, 11)

📄 Column names: ['num_preg', 'glucose_conc', 'diastolic_bp', 'thickness', 'insulin', 'bmi', 'diab_pred', 'age', 'skin', 'diabetes', 'age_group']

📊 Data types:

```

num_preg      int64
glucose_conc  int64
diastolic_bp  int64
thickness     int64
insulin       int64
bmi           float64
diab_pred     float64
age           int64
skin          float64
diabetes      bool
age_group     category
dtype: object

```

```
In [48]: print("\n ? Missing values:\n", df.isnull().sum())
```

? Missing values:

```

num_preg      0
glucose_conc  0
diastolic_bp  0
thickness     0
insulin       0
bmi           0
diab_pred     0
age           0
skin          0
diabetes      0
age_group     0
dtype: int64

```


4. CLEANING AND PREPROCESS

```
In [52]: #Convert diabetes to 0/1:
df['diabetes'] = df['diabetes'].astype(int)
```

```
In [55]: #Replace invalid 0s with median:
invalid_zero_cols = ['glucose_conc', 'diastolic_bp', 'thickness', 'insulin', 'bmi']
for col in invalid_zero_cols:
    median = df[col].median()
    df[col] = df[col].replace(0, median)
```


5. DATA VISUALISATIONS AND EDA

```
In [56]: #Class Distribution
sns.countplot(data=df, x='diabetes')
plt.title("Distribution of Diabetes Cases (0 = No, 1 = Yes)")
plt.xlabel("Diabetes Diagnosis")
plt.ylabel("Number of Patients")
plt.show()
```



```
In [64]: #Correlation Heatmaps
# Select only numeric features for correlation
numeric_df = df.select_dtypes(include=['int64', 'float64'])

# Plot the heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(numeric_df.corr(), annot=True, cmap='coolwarm', fmt=".2f")
plt.title("Feature Correlation Heatmap")
plt.show()
```



```
In [ ]: #Box Plots
features = df.columns.drop('diabetes')
for col in features:
    plt.figure(figsize=(6, 4))
    sns.boxplot(data=df, x='diabetes', y=col)
    plt.title(f"{col} vs. Diabetes")
    plt.show()
```

6. Predictive Modeling – Logistic Regression

```
In [58]: #Train Test Split
X = df.drop('diabetes', axis=1)
y = df['diabetes']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42, stratify=y)
```

```
In [67]: # Train and evaluate the model
#Define features and target
X = df.drop(columns=['diabetes', 'age_group', 'checkup_date']) # drop non-numeric or derived columns
y = df['diabetes']

#Train-test split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y)

#Train the model
```

```

model = LogisticRegression(max_iter=1000)
model.fit(X_train, y_train)

#Predict and evaluate
y_pred = model.predict(X_test)
print("🔗 Accuracy:", accuracy_score(y_test, y_pred))
print("\n📊 Classification Report:\n", classification_report(y_test, y_pred))
print("\n📊 Confusion Matrix:\n", confusion_matrix(y_test, y_pred))

```

🔗 Accuracy: 0.7207792207792207

📊 Classification Report:

	precision	recall	f1-score	support
0	0.77	0.82	0.79	100
1	0.62	0.54	0.57	54
accuracy			0.72	154
macro avg	0.69	0.68	0.68	154
weighted avg	0.71	0.72	0.72	154

📊 Confusion Matrix:

```

[[82 18]
 [25 29]]

```

7. STATISTICAL TEST EXAMPLE

```

In [68]: group_0 = df[df['diabetes'] == 0]['glucose_conc']
group_1 = df[df['diabetes'] == 1]['glucose_conc']

stat, p = ttest_ind(group_0, group_1)
print(f"t-statistic: {stat:.3f}, p-value: {p:.4f}")

```

t-statistic: -15.674, p-value: 0.0000

8. DATA MANIPULATION

```

In [69]: glucose_25 = np.percentile(df['glucose_conc'], 25)
glucose_75 = np.percentile(df['glucose_conc'], 75)
print(f"Glucose IQR: {glucose_25}-{glucose_75}")

df['high_glucose'] = np.where(df['glucose_conc'] > glucose_75, 1, 0)

```

Glucose IQR: 99.75-140.25

```

In [71]: df['age_group'] = pd.cut(df['age'], bins=[20, 30, 40, 50, 60, 100],
                                labels=['20s', '30s', '40s', '50s', '60+'])
print(df.groupby('age_group', observed=True)['diabetes'].mean())
print(df.groupby('age_group', observed=False)['diabetes'].mean())

```

```


age_group
20s    0.215827
30s    0.484076
40s    0.566372
50s    0.574074
60+    0.259259
Name: diabetes, dtype: float64
age_group
20s    0.215827
30s    0.484076
40s    0.566372
50s    0.574074
60+    0.259259
Name: diabetes, dtype: float64

```

9. TIME SIMULATION + LINE PLOT

```
In [72]: date_range = pd.date_range(start='2022-01-01', periods=len(df), freq='D')
df['checkup_date'] = date_range

diabetes_ts = df.groupby(df['checkup_date'].dt.to_period("M"))['diabetes'].mean()
diabetes_ts.plot(title="Simulated Monthly Diabetes Rate")
plt.show()
```



10. PROBABILISTIC PREDICTION OUTPUT (SOLUTION)

```
In [73]: probs = model.predict_proba(X_test)
print(probs[:5]) # class 0 vs class 1 probabilities for 5 samples

[[0.3850491  0.6149509 ]
 [0.91034925 0.08965075]
 [0.70847737 0.29152263]
 [0.7374493  0.2625507 ]
 [0.97233773 0.02766227]]
```